

Your Paper Title

Author One^{1,†,*} Author Two^{1,2,†} Author Three²

¹ Institution A ² Institution B

author_one@example.org author_two@example.org

[†] Equal contribution. ^{*} Corresponding author.

Abstract

Replace this abstract with a concise account of the research question, approach, main findings, and limitations. Explain what is evaluated and why the evidence supports the central claim. Define any essential terms and keep the abstract self-contained. This sample demonstrates the Agent Evaluation Science template, including author information, citations, figures, tables, equations, footnotes, and appendices. All numerical results below are invented for illustration.

1 Introduction

Introduce the research problem, explain its significance, and state the paper's contribution. Identify the system, task, or evaluation resource under study. Describe the gap addressed by the work and distinguish it from the closest existing approaches. The text in this document is a writing scaffold, not a report of completed research.

A new paragraph has no first-line indentation. The template supplies paragraph spacing, heading sizes, and page margins. Use ordinary LaTeX markup such as **bold**, *italics*, and <https://evalscience.org/>. Keep formatting commands out of prose unless they communicate a meaningful distinction. A footnote can hold a brief supporting detail.¹

2 Related Work

A narrative citation uses [Kiela et al. \(2021\)](#). A parenthetical citation uses ([Rajpurkar et al., 2016](#)). Multiple sources can share a citation ([Chiang et al., 2024](#); [Geiger et al., 2012](#)). These real publications provide citation examples; replace them with sources relevant to your paper.

Explain how prior approaches define their tasks, gather evidence, and assess outcomes. Connect each cited source to a specific statement. When revisiting a source, cite it again where needed ([Kiela et al., 2021](#)). The numbers after "Cited at:" in References return to individual citation occurrences, including repeated citations on the same page; they are not page numbers.

3 Method

Describe the procedure in enough detail that another researcher can understand what was measured and reproduce the essential steps. Separate decisions about task construction, experimental conditions, and outcome assessment. State which parts of the procedure were fixed before examining the results.

3.1 Evaluation Setup

Specify the inputs, available tools, environment, and resource limits. For example, an evaluation description can identify:

¹Use footnotes for concise qualifications or supplementary information. Keep essential evidence in the main text.

- The tasks and the criteria used to select them.
- The information and tools available to each system.
- The conditions under which an attempt is considered complete.

Record relevant model versions and evaluation dates. Explain whether attempts are independent and how failures, retries, and incomplete records are handled.

3.1.1 Outcome Measures

Let N be the number of evaluated attempts and $s_i \in \{0, 1\}$ indicate whether attempt i meets the stated success criteria. The completion rate is

$$\hat{S} = \frac{1}{N} \sum_{i=1}^N s_i. \quad (1)$$

Define the unit of analysis and the denominator for every measure. If partial success or abstention is possible, explain how it is represented rather than silently removing those outcomes from Equation (1).

3.2 Procedure

A numbered list is useful when order matters:

1. Initialize the environment and provide the task instruction.
2. Record actions, observations, and resource use during execution.
3. Assess the outcome using the predefined scoring procedure.

Figure 1 summarizes the relationship between these stages. Use descriptive labels and refer to figures and tables by number.



Figure 1 | An illustrative evaluation workflow. Replace this diagram and caption with a figure that supports the paper’s argument.

4 Results and Discussion

Present the main findings and explain what they establish. Table 1 uses invented values to demonstrate a compact results table. State the meaning of each column, the direction of improvement, and the evaluation conditions. Use actual observations and suitable uncertainty estimates in a research paper.

System	Completion (%)	Mean actions
Baseline	61.0	14.2
Variant A	68.5	12.8
Variant B	72.0	13.1

Table 1 | Synthetic results for layout demonstration. Completion is the percentage of successful attempts; mean actions measures resource use.

Discuss both successful and unsuccessful cases. A higher aggregate score may hide differences in task coverage, failure modes, or resource use. Explain the scope of the comparison

and connect quantitative results to the evidence needed to interpret them. Appendix A illustrates supplementary material.

5 Limitations

Describe which settings, populations, or tasks the evidence does not cover. Explain limitations of the measurement procedure and any unresolved sources of uncertainty. The limitations of this sample are straightforward: no experiment was run, and the results table contains fictional values.

6 Conclusion

State what the work establishes and identify the most consequential remaining question. Keep the conclusion consistent with the reported evidence. Remove all sample prose, fictional results, and unused examples before sharing your paper.

Acknowledgments

Replace with acknowledgments, or delete this block.

Ethics Statement

Describe relevant ethical considerations, or delete this section if not applicable.

References

- Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, et al. 2024. Chatbot arena: An open platform for evaluating LLMs by human preference. *arXiv preprint arXiv:2403.04132*. <https://doi.org/10.48550/arXiv.2403.04132>. Cited at: 1.
- David Geiger, Michael Rosemann, Erwin Fieft, and Martin Schader. Crowdsourcing information systems: Definition, typology, and design. In *Proceedings of the 33rd International Conference on Information Systems*, 2012. <https://madoc.bib.uni-mannheim.de/32631/>. Cited at: 1.
- Douwe Kiela, Max Bartolo, Yixin Nie, Divyansh Kaushik, Atticus Geiger, Zhengxuan Wu, et al. Dynabench: Rethinking benchmarking in NLP. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4110–4124. Association for Computational Linguistics, 2021. <https://doi.org/10.18653/v1/2021.naacl-main.324>. Cited at: 1, 2.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392. Association for Computational Linguistics, 2016. <https://doi.org/10.18653/v1/D16-1264>. Cited at: 1.

A Additional Evaluation Details

Place supporting protocols, additional analyses, and extended results here. Appendix figures and tables use labels such as Figure A1 and Table A1. Refer to the relevant appendix from the main text so that readers know why the additional material matters.

A.1 Reporting Details

- Provide the task selection criteria and any exclusions.
- Record model identifiers, tool access, and resource budgets.
- Describe scoring, review, and disagreement resolution.

Field	Description
Task ID	Stable identifier for the evaluated task.
Outcome	Success, failure, or an explicitly defined unresolved state.
Evidence	The observation or record supporting the assigned outcome.

Table A1 | An example reporting schema for supplementary evaluation records.